



# Process-First AI:

A Practical Framework for MSME Adoption

**Why 74% of AI Initiatives Fail—  
And How Mid-Size Enterprises Can Beat the Odds**

By Tushar Gandhi & Omkar Palit

**December 2025**

## Table of Contents

Table of Contents	2
Executive Summary	3
<b>1. The AI Adoption Crisis .....</b>	<b>4</b>
1.1 The Platform-FirstTrap	4
1.2 The MSME Challenge	5
<b>2. The Process-First Framework .....</b>	<b>6</b>
2.1 Learning from High Performers	6
2.2 The Five-Step Methodology	6
Step 1: Workflow Audit	7
Step 2: Impact Quantification	7
Step 3: AI Suitability Assessment	7
Step 4: Build a Focused Solution	8
Step 5: Measure and Iterate	8
<b>3. Proof of Concept: Status Synthesiser .....</b>	<b>9</b>
3.1 Problem and Impact	9
3.2 Suitability and Solution	9
3.3 Business Impact and ROIAnalysis	9
<b>4. Implementation Considerations .....</b>	<b>15</b>
4.1 Build vs. Buy	15
4.2 Technology Stack	15
4.3 Measuring Success	15
<b>5. Conclusion .....</b>	<b>16</b>
<b>References</b>	17
<b>About the Authors</b>	18

## Executive Summary

Despite 88% of organisations using AI, 74% of initiatives fail to demonstrate tangible value, and fewer than 20% ever scale enterprise-wide. Mid-size enterprises face an impossible choice: pay premium prices for enterprise platforms or risk falling behind. Neither addresses the fundamental problem: starting with technology instead of business problems.

The Process-First Framework inverts this sequence. By starting with workflow audits, quantifying impact, applying rigorous suitability assessment, building focused solutions, and measuring results, organisations achieve transformative outcomes. High performers are 3.6× more likely to pursue transformative goals and achieve 2× the ROI by fundamentally redesigning workflows rather than automating existing processes.

Our POC, Status Synthesiser, demonstrates this framework in practice: we quantified a \$156,000 annual opportunity, scored 7/7 on the AI Suitability Assessment, and built a working proof-of-concept in 25 hours for \$2,400. The POC successfully automates 90% of the manual effort. Depending on deployment approach—public LLM APIs or self-hosted models—projected Year 1 ROI ranges from 1,700% to 2,648%, with ongoing returns of 2,456% to 5,084% annually. Both scenarios include managed service support, ensuring production reliability.

Success depends 70% on people and process integration, only 30% on technology. Mid-size enterprises that master this methodology achieve enterprise-grade AI capabilities without enterprise-grade budgets.

## 1. The AI Adoption Crisis

Corporate AI spending has reached unprecedented levels, yet the vast majority of investments fail to deliver measurable value. McKinsey's November 2025 State of AI report reveals a stark progression: while 88% of organisations use AI in at least one function, 62% remain stuck in pilot phases, 74% fail to show tangible value, only 25% achieve expected ROI, and fewer than 20% scale enterprise-wide. MIT's August 2025 research quantifies this starkly: 95% of enterprise AI pilots fail to achieve measurable impact.

Perhaps most telling is continued optimism despite failures. DeepQuery's June 2025 survey found 60% of organisations plan to increase AI spending in 2026, even as current investments underdeliver. This disconnect suggests organisations recognise AI's potential but lack effective frameworks for realising it.

### THE AI ADOPTION GAP

The Narrowing Funnel of AI Success



#### 1.1 The Platform-First Trap

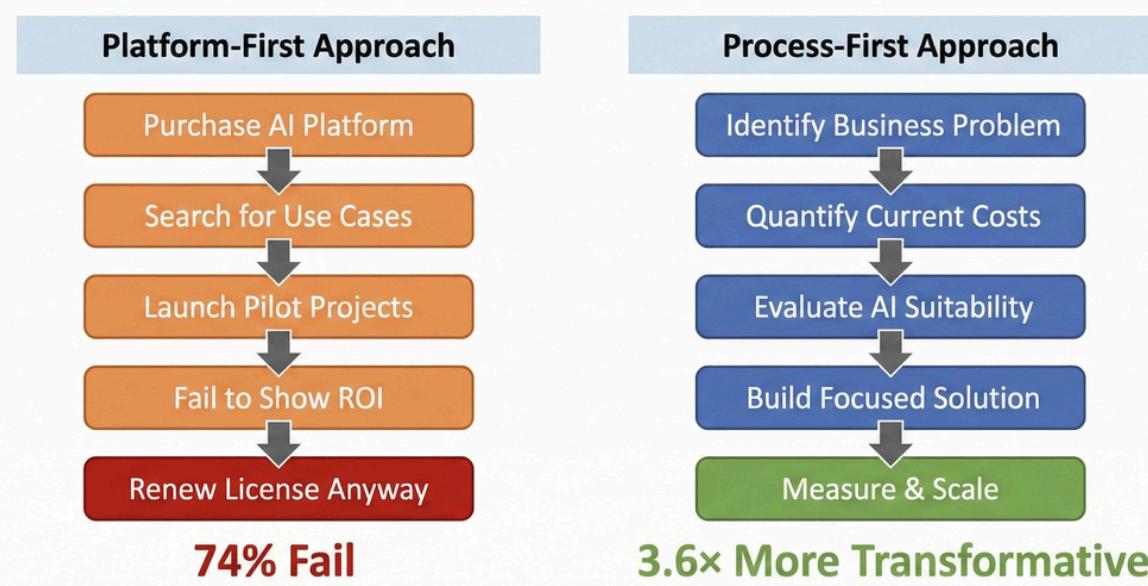
- The root cause follows a predictable pattern:
  - Purchase an enterprise AI platform
  - Struggle to identify use cases
  - Launch pilots without clear metrics
  - Fail to demonstrate ROI
  - Renew license citing strategic importance.

This approach fundamentally misunderstands how AI creates value. By starting with technology and searching for problems, organisations optimise for vendor satisfaction rather than business outcomes.

- The Process-First alternative inverts this:
  - Identify a high-impact business problem
  - Quantify current costs and establish success metrics
  - Evaluate if AI is appropriate
  - Build a focused tool with a clear ROI target
  - Measure results and scale or pivot based on data.

McKinsey's 2025 research validates this. Organisations that fundamentally redesign workflows are 2.8× more likely to succeed, high performers pursue transformative goals at 3.6× the rate of typical organisations, and they achieve 2× the ROI by focusing on half as many opportunities.

## PLATFORM-FIRST VS PROCESS-FIRST



### 1.2 The MSME Challenge

For mid-size enterprises with 100-5,000 employees, the Platform-First trap creates unique challenges. Enterprise AI platforms typically cost \$100,000-\$500,000 annually, yet are optimised for Fortune 500 scale. While major LLM providers maintain GDPR and HIPAA compliance, many organisations face internal policies requiring data residency, compliance departments unfamiliar with API-level controls, or board-level risk aversion to external AI services. Additionally, mid-size companies rarely have dedicated AI teams, and technology leaders manage daily operations while driving strategic initiatives. The opportunity cost of failed pilots is measured not just in dollars, but in leadership credibility and organisational momentum.

BCG's 10-20-70 rule provides crucial context: AI success depends 10% on algorithms, 20% on technology infrastructure, and 70% on people and process integration. The Process-First Framework addresses these challenges by starting with the 70%, not the 30%.

The conventional approach asks, “We have AI, now what?” The right question for mid-size enterprises is “*We have problem X—can AI help?*” Starting with workflow deep-dives rather than platform capabilities inverts the failure pattern.

## 2. The Process-First Framework

The Process-First Framework inverts the conventional AI adoption sequence by starting with business problems rather than technology platforms. Instead of purchasing AI capabilities and searching for applications, organisations first identify high-impact workflows, quantify their costs, assess AI suitability, build focused solutions, and measure results against explicit criteria.

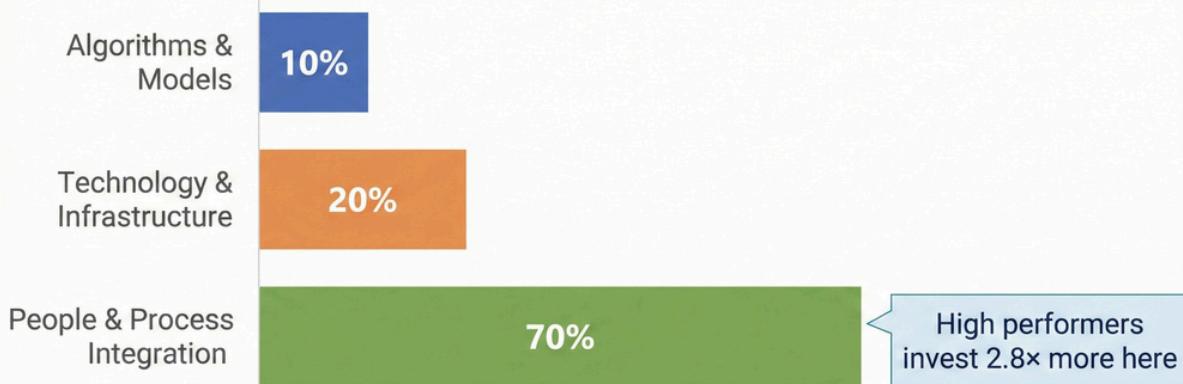
This methodology addresses the fundamental disconnect revealed in research: 74% of initiatives fail because organisations optimise for vendor satisfaction rather than business outcomes. By anchoring every decision in workflow economics and measurable impact, the framework ensures technology serves business needs rather than the reverse. The framework consists of two foundational elements: understanding what high-performing organisations do differently (Section 2.1) and applying a disciplined five-step methodology to identify and validate opportunities (Section 2.2).

### 2.1 Learning from High Performers

McKinsey's November 2025 analysis reveals high performers share distinct characteristics: they pursue transformative goals at 3.6× the rate of typical organisations, are 2.8× more likely to fundamentally redesign workflows, commit 20%+ of digital budgets to AI while focusing on half as many opportunities, and invest heavily in change management and capability building. BCG's 10-20-70 rule quantifies this: 70% of success depends on people and process integration, with only 20% on technology and 10% on algorithms. BCG's 10-20-70 rule provides crucial context: AI success depends 10% on algorithms, 20% on technology infrastructure, and 70% on people and process integration. The Process-First Framework addresses these challenges by starting with the 70%, not the 30%.

# HIGH PERFORMERS' RESOURCE ALLOCATION

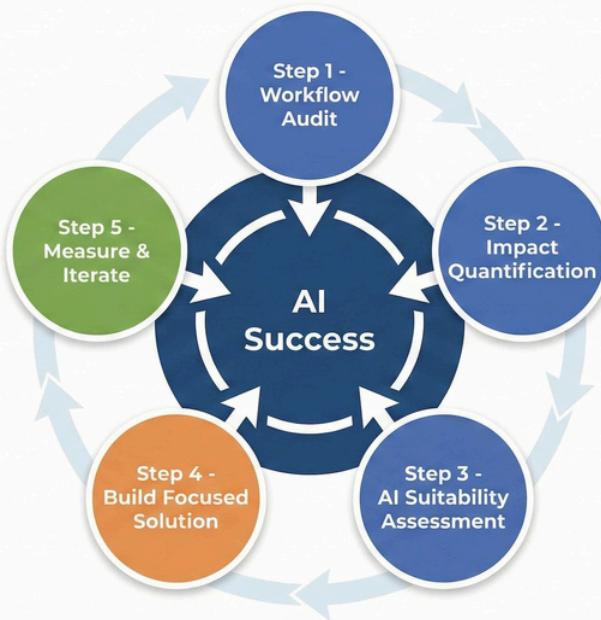
BCG 10-20-70 Rule: Where AI Success Really Comes From



Success depends 70% on people and process, only 30% on technology

## 2.2 The Five-Step Methodology

### THE PROCESS-FIRST FRAMEWORK



#### Step 1: Workflow Audit

Map existing business processes before evaluating technology. Focus on three categories:

- High-Volume Tasks: Daily/weekly recurring activities performed by multiple team members
- Time-Intensive Processes: Tasks requiring 2+ hours per occurrence
- Quality Issues: Processes with inconsistent output quality by person or circumstance

## Step 2: Impact Quantification

For top candidates, calculate comprehensive impact: (1) Time Costs: [Hours per occurrence] × [Annual occurrences] × [Number of people], (2) Financial Impact: Convert time to dollars using fully loaded rates, (3) Quality Costs: Estimate costs of missed risks, errors, and rework, (4) Opportunity Costs: Consider what team members could accomplish with recovered time. This quantification establishes baseline metrics and creates a prioritised opportunity list. Processes with \$100,000+ annual impact warrant focused attention.

## Step 3: AI Suitability Assessment

Not every high-impact process benefits from AI. The AI Suitability Scorecard evaluates seven criteria, each scored 0 (Poor Fit), 0.5 (Moderate), or 1.0 (Strong Fit): Volume/Frequency, Manual Effort, Cognitive Load, Consistency Need, Data Availability, Error Tolerance, and Pain Level. Score 5.0-7.0 indicates excellent AI candidate; 3.0-4.5 suggests potential with modifications; 0.0-2.5 recommends alternative solutions.

## AI SUITABILITY SCORECARD

### 7 Criteria for Evaluating AI Opportunities

CRITERION	STRONG FIT:	POOR FIT:	SCORE
Volume/Frequency	Daily/Weekly	Quarterly+	● ○ ○
Manual Effort	2+ hours	<15 minutes	● ○ ○
Cognitive Load	Pattern recognition	Creative judgment	● ○ ○
Consistency Need	Standard format	Highly contextual	○ ○ ○
Data Availability	Text accessible	External APIs	● ○ ○
Error Tolerance	Human review OK	Zero errors	● ○ ○
Pain Level	Frequent complaints	Works fine	● ○ ○

● = Strong Fit (1.0)

○ = Moderate Fit (0.5)

● = Poor Fit (0.0)

**Score 5-7:** Excellent AI Candidate | **3-4:** Moderate Fit | **0-2:** Consider Alternatives

## Step 4: Build a Focused Solution

For processes scoring 5.0+, build a proof of concept with: (1) Single-purpose design solving one problem exceptionally well, (2) Clear success metrics defined before coding (80%+ time reduction, 85%+ accuracy, 70%+ adoption, NPS +50), (3) Standard frameworks (LangChain, FastAPI, React), (4) 25-30 hour timeline forcing ruthless prioritization, (5) Real data testing from day one.

## Step 5: Measure and Iterate

Establish measurement cadence: Week 1 (usability and adoption), Week 4 (adoption trajectory and impact signals), Week 8 (comprehensive evaluation and scale-or-pivot decision), Quarterly (sustained impact and enhancement opportunities). Decision framework at Week 8: Scale (adoption >70%, metrics

exceeding targets), Iterate (mixed results with clear improvement path), Pivot (low adoption but alternative identified), Abandon (fundamental flaws or better opportunities).

Platform vendors optimise for breadth. Mid-size enterprises must optimise for depth. Focusing on three well-understood, quantified workflows delivers better outcomes than fifty mediocre automations.

### **3. Proof of Concept: Status Synthesiser**

#### **3.1 Problem and Impact**

Project managers at a mid-size technology services company spent 3-5 hours weekly compiling status reports from Slack, email, and Jira. For 10 PMs, this meant 2,340 annual hours at \$75/hour = \$175,500 in direct costs. Post-mortem analysis revealed early warning signals existed 2-3 weeks before escalation in 30% of delayed projects, adding \$40,000-\$60,000 in quality costs. Opportunity costs from forgone strategic work added another \$30,000-\$60,000 annually. Total quantified impact: \$245,000-\$295,000.

#### **3.2 Suitability and Solution**

The AI Suitability Assessment scored 7.0/7.0: high volume (52 reports/year/PM), significant manual effort (3-5 hours), pattern recognition task, standard format required, text inputs accessible, human review acceptable, and consistent complaints. This perfect score indicated strong alignment between problem characteristics and AI capabilities.

The solution implements a multi-agent architecture: Agent 1 (Extraction) processes unstructured communications and detects scope creep, Agent 2 (Analysis) applies risk scoring and delta analysis comparing current versus previous reports, and Agent 3 (Synthesis) generates executive-ready Markdown reports. Technology stack: Gemini 2.5 Flash (adaptable to Llama 3.1), LangChain for orchestration, Pydantic for structured output, FastAPI backend, React frontend. Development: 25 hours over six days.

#### **3.3 Business Impact and ROI Analysis**

Applying the framework to status report generation reveals quantifiable impact:

### **Problem Quantification:**

- Current state:  $10 \text{ PMs} \times 4 \text{ hours/week} \times 52 \text{ weeks} = 2,080 \text{ hours annually}$
- Direct cost:  $2,080 \text{ hours} \times \$75/\text{hour} = \$156,000/\text{year}$
- Quality costs:  $\$40,000-\$50,000$  (from delayed risk detection in post-mortems)
- Total opportunity:  $\$200,000-\$250,000$  annually

### **Solution Development:**

- POC built in 25 hours using Gemini 2.5 Flash, LangChain, FastAPI, React
- Development investment: \$2,400 (labor + tooling with opportunity cost)
- Technical validation: Multi-agent system successfully automates ~90% of manual tasks
- Conservative time reduction assumption: 85%
- Projected annual savings: \$132,600

A focused 25-hour build delivered 1,700-2,648% Year 1 ROI—proving mid-size enterprises don't need enterprise budgets or specific technology stacks. Both public APIs and self-hosted deployments deliver exceptional returns when guided by disciplined methodology. The barrier isn't technology—it's process rigour.

### **Total Cost of Ownership: Two Deployment Options**

Organisations can choose between public LLM APIs or self-hosted deployment based on data governance requirements. Both scenarios include managed service support (\$200/month) covering monitoring, updates, bug fixes, and enhancements.

## **4. Implementation Considerations**

### **4.1 Build vs. Buy**

Build when: unique workflows, data governance requires local deployment, cost constraints, deep integration needs, or competitive advantage. Buy when: commodity use cases, enterprise support required, rapid deployment critical, limited technical capacity, or broad feature requirements (10+ distinct capabilities). For mid-size enterprises, building focused tools often delivers superior ROI for initial use cases.

## **4.2 Technology Stack**

LLM Selection: API-based (Gemini 2.5 Flash, Claude) offers immediate setup, \$0.03-0.10 per 1K tokens, cutting-edge performance, and zero maintenance. Local deployment (Llama, Mistral) requires 2-4 hours of setup but provides complete data control at zero ongoing cost beyond compute.

Recommendation: Start with the API for POC speed, evaluate locally for production if governance requires. Development: Use LangChain/LangGraph for orchestration, Pydantic for structured output, FastAPI for backend, React for UI.

## **4.3 Measuring Success**

Track financial metrics (cost avoidance, revenue impact, quality costs), operational metrics (adoption rate target 70%+, time savings target 80%+, accuracy target 85%+), and user satisfaction (NPS target +50, task ease 8+/10, output trust 70%+). Measure at Week 1, Week 4, Week 8, then quarterly.

## **5. Conclusion**

The AI adoption crisis facing mid-size enterprises stems from fundamental misalignment between how AI is sold and how it creates value. The Process-First Framework inverts this by starting with workflow audits, quantifying impact, applying rigorous suitability assessment, building focused solutions, and measuring results against explicit criteria.

The Status Synthesiser demonstrates this framework in practice: \$2,400 investment in a 25-hour POC targeting a \$156,000 annual problem, with projected ROI of 1,700-2,648% depending on deployment approach. The framework applies equally well to other high-volume cognitive tasks: meeting notes to action items, change request auditing, and proactive risk detection all score 6.5-7.0 on the AI Suitability Assessment.

Success does not require enterprise budgets or comprehensive platforms. It requires starting with problems rather than platforms, applying disciplined methodology, and measuring honestly. The tools and frameworks exist. The opportunity is clear. The choice is whether to follow a proven path or repeat the mistakes that leave 74% of initiatives without demonstrable value.

## References

- Boston Consulting Group. (2024, October). Ten Commandments of Applied AI: What the 10% of Companies Successfully Scaling AI Are Doing Differently. *Boston Consulting Group*.
- DeepQuery. (2025, June). AI Investment Trends Survey 2025: Enterprise Spending Patterns and Value Realisation. *DeepQuery Research*.
- Fullview. (2025). AI Project Success Rates Report: Understanding Why 70-85% of AI Initiatives Fail. *Fullview Analytics*.
- McKinsey & Company. (2025, November). The State of AI in 2025: High Performers Set Themselves Apart. *McKinsey Global Institute*.
- MIT NANDA. (2025, August). Enterprise AI Pilot Failure Rates: A Comprehensive Analysis of the 95% Failure Rate in Achieving Measurable Impact. *MIT National AI and Digital Advancement Centre*.
- Stack AI. (2025). AI ROI and Scaling Statistics: Why Only 25% Achieve Expected Returns and Fewer Than 20% Scale Enterprise-Wide. *Stack AI Research*.

## About the Authors

### **Omkar Palit**

*Omkar brings 15 years of experience in program management, business development, enterprise sales, and business excellence. His career spans leading digital transformation initiatives for mid-size enterprises and optimising complex operational processes for measurable business impact. He has diverse experience across oil and gas, energy and digital SaaS industries.*

### **Tushar Gandhi**

*Tushar is a Cloud & AI Consultant with 14+ years delivering enterprise technology solutions and 12+ months specialised in AI/GenAI. Multi-certified across AWS, Azure, and GCP with proven expertise in cloud migration strategy and GenAI implementation. Track record includes cloud assessment, leading digital transformation projects, developing AI-powered automation solutions, pre-sales consulting and 12 months of international experience.*

**[Connect on LinkedIn to continue the conversation and access additional resources.](#)**